

# Bark 🐶

[↑ Back to top](#)

Bark is a multi-lingual TTS model created by [Suno-AI](#). It can generate conversational speech as well as music and sound effects. It is architecturally very similar to Google's [AudioLM](#). For more information, please refer to the [Suno-AI's repo](#).

## Acknowledgements

- 🏰 [Suno-AI](#) for training and open-sourcing this model.
- 🏰 [gitmylo](#) for finding [the solution](#) to the semantic token generation for voice clones and finetunes.
- 🏰 [serp-ai](#) for controlled voice cloning.

## Example Use

```
text = "Hello, my name is Manmay , how are you?"

from TTS.tts.configs.bark_config import BarkConfig
from TTS.tts.models.bark import Bark

config = BarkConfig()
model = Bark.init_from_config(config)
model.load_checkpoint(config, checkpoint_dir="path/to/model/dir/", eval=True)

# with random speaker
output_dict = model.synthesize(text, config, speaker_id="random", voice_dirs=None)

# cloning a speaker.
# It assumes that you have a speaker file in `bark_voices/speaker_n/speaker.wav` or `bark_voices/...`
output_dict = model.synthesize(text, config, speaker_id="ljspeech", voice_dirs="bark_voices/")
```

Using 🐸 TTS API:

[v: dev](#) ▼

```

from TTS.api import TTS

# Load the model to GPU
# Bark is really slow on CPU, so we recommend using GPU.
tts = TTS("tts_models/multilingual/multi-dataset/bark", gpu=True)

# Cloning a new speaker
# This expects to find a mp3 or wav file like `bark_voices/new_speaker/speaker.wav`
# It computes the cloning values and stores in `bark_voices/new_speaker/speaker.npz`
tts.tts_to_file(text="Hello, my name is Manmay , how are you?",
                file_path="output.wav",
                voice_dir="bark_voices/",
                speaker="ljspeech")

# When you run it again it uses the stored values to generate the voice.
tts.tts_to_file(text="Hello, my name is Manmay , how are you?",
                file_path="output.wav",
                voice_dir="bark_voices/",
                speaker="ljspeech")

# random speaker
tts = TTS("tts_models/multilingual/multi-dataset/bark", gpu=True)
tts.tts_to_file("hello world", file_path="out.wav")

```

↑ Back to top

Using 🦊 TTS Command line:

```

# cloning the `ljspeech` voice
tts --model_name tts_models/multilingual/multi-dataset/bark \
--text "This is an example." \
--out_path "output.wav" \
--voice_dir bark_voices/ \
--speaker_idx "ljspeech" \
--progress_bar True

# Random voice generation
tts --model_name tts_models/multilingual/multi-dataset/bark \
--text "This is an example." \
--out_path "output.wav" \
--progress_bar True

```

v: dev ▼

# Important resources & papers

[↑ Back to top](#)

- Original Repo: <https://github.com/suno-ai/bark>
- Cloning implementation: <https://github.com/serp-ai/bark-with-voice-clone>
- AudioLM: <https://arxiv.org/abs/2209.03143>

# BarkConfig

[↑ Back to top](#)

```
class TTS.tts.configs.bark_config.BarkConfig(output_path='output', logger_uri=None,
run_name='run', project_name=None, run_description='🐶 Coqui trainer run.',
print_step=25, plot_step=100, model_param_stats=False, wandb_entity=None,
dashboard_logger='tensorboard', save_on_interrupt=True, log_model_step=None,
save_step=10000, save_n_checkpoints=5, save_checkpoints=True, save_all_best=False,
save_best_after=10000, target_loss=None, print_eval=False, test_delay_epochs=0,
run_eval=True, run_eval_steps=None, distributed_backend='nccl',
distributed_url='tcp://localhost:54321', mixed_precision=False, precision='fp16',
epochs=1000, batch_size=32, eval_batch_size=16, grad_clip=0.0,
scheduler_after_epoch=True, lr=0.001, optimizer='radam', optimizer_params=None,
lr_scheduler=None, lr_scheduler_params=<factory>, use_grad_scaler=False,
allow_tf32=False, cudnn_enable=True, cudnn_deterministic=False, cudnn_benchmark=False,
training_seed=54321, model='bark', num_loader_workers=0, num_eval_loader_workers=0,
use_noise_augment=False, audio=<factory>, use_phonemes=False, phonemizer=None,
phoneme_language=None, compute_input_seq_cache=False, text_cleaner=None,
enable_eos_bos_chars=False, test_sentences_file='', phoneme_cache_path=None,
characters=None, add_blank=False, batch_group_size=0, loss_masking=None,
min_audio_len=1, max_audio_len=inf, min_text_len=1, max_text_len=inf,
compute_f0=False, compute_energy=False, compute_linear_spec=False,
precompute_num_workers=0, start_by_longest=False, shuffle=False, drop_last=False,
datasets=<factory>, test_sentences=<factory>, eval_split_max_size=None,
eval_split_size=0.01, use_speaker_weighted_sampler=False,
speaker_weighted_sampler_alpha=1.0, use_language_weighted_sampler=False,
language_weighted_sampler_alpha=1.0, use_length_weighted_sampler=False,
length_weighted_sampler_alpha=1.0, num_chars=0, semantic_config=<factory>,
fine_config=<factory>, coarse_config=<factory>, CONTEXT_WINDOW_SIZE=1024,
SEMANTIC_RATE_HZ=49.9, SEMANTIC_VOCAB_SIZE=10000, CODEBOOK_SIZE=1024,
N_COARSE_CODEBOOKS=2, N_FINE_CODEBOOKS=8, COARSE_RATE_HZ=75, SAMPLE_RATE=24000,
USE_SMALLER_MODELS=False, TEXT_ENCODING_OFFSET=10048, SEMANTIC_PAD_TOKEN=10000,
TEXT_PAD_TOKEN=129595, SEMANTIC_INFER_TOKEN=129599, COARSE_SEMANTIC_PAD_TOKEN=12048,
COARSE_INFER_TOKEN=12050, REMOTE_MODEL_PATHS=None, LOCAL_MODEL_PATHS=None,
SMALL_REMOTE_MODEL_PATHS=None, CACHE_DIR='/home/docs/.local/share/tts/suno/bark_v0',
DEF_SPEAKER_DIR='/home/docs/.local/share/tts/bark_v0/speakers')
```

[\[source\]](#)

 v: dev ▼

## Bark TTS configuration

## PARAMETERS:

- **model** (*str*) – model name that registers the model. ↑ Back to top
- **audio** (*BarkAudioConfig*) – audio configuration. Defaults to `BarkAudioConfig()`.
- **num\_chars** (*int*) – number of characters in the alphabet. Defaults to 0.
- **semantic\_config** (*GPTConfig*) – semantic configuration. Defaults to `GPTConfig()`.
- **fine\_config** (*FineGPTConfig*) – fine configuration. Defaults to `FineGPTConfig()`.
- **coarse\_config** (*GPTConfig*) – coarse configuration. Defaults to `GPTConfig()`.
- **CONTEXT\_WINDOW\_SIZE** (*int*) – GPT context window size. Defaults to 1024.
- **SEMANTIC\_RATE\_HZ** (*float*) – semantic tokens rate in Hz. Defaults to 49.9.
- **SEMANTIC\_VOCAB\_SIZE** (*int*) – semantic vocabulary size. Defaults to 10\_000.
- **CODEBOOK\_SIZE** (*int*) – codec codebook size. Defaults to 1024.
- **N\_COARSE\_CODEBOOKS** (*int*) – number of coarse codebooks. Defaults to 2.
- **N\_FINE\_CODEBOOKS** (*int*) – number of fine codebooks. Defaults to 8.
- **COARSE\_RATE\_HZ** (*int*) – coarse tokens rate in Hz. Defaults to 75.
- **SAMPLE\_RATE** (*int*) – sample rate. Defaults to 24\_000.
- **USE\_SMALLER\_MODELS** (*bool*) – use smaller models. Defaults to False.
- **TEXT\_ENCODING\_OFFSET** (*int*) – text encoding offset. Defaults to 10\_048.
- **SEMANTIC\_PAD\_TOKEN** (*int*) – semantic pad token. Defaults to 10\_000.
- **TEXT\_PAD\_TOKEN** (*[type]*) – text pad token. Defaults to 10\_048.
- **TEXT\_EOS\_TOKEN** (*[type]*) – text end of sentence token. Defaults to 10\_049.
- **TEXT\_SOS\_TOKEN** (*[type]*) – text start of sentence token. Defaults to 10\_050.
- **SEMANTIC\_INFER\_TOKEN** (*int*) – semantic infer token. Defaults to 10\_051.
- **COARSE\_SEMANTIC\_PAD\_TOKEN** (*int*) – coarse semantic pad token. Defaults to 12\_048.
- **COARSE\_INFER\_TOKEN** (*int*) – coarse infer token. Defaults to 12\_050.
- **REMOTE\_BASE\_URL** (*[type]*) – remote base url. Defaults to ["https://huggingface.co/erogol/bark/tree"](https://huggingface.co/erogol/bark/tree).
- **REMOTE\_MODEL\_PATHS** (*Dict*) – remote model paths. Defaults to None.
- **LOCAL\_MODEL\_PATHS** (*Dict*) – local model paths. Defaults to None.
- **SMALL\_REMOTE\_MODEL\_PATHS** (*Dict*) – small remote model paths. Defaults to None.
- **CACHE\_DIR** (*str*) – local cache directory. Defaults to `get_user_data_dir()`.
- **DEF\_SPEAKER\_DIR** (*str*) – default speaker directory to store speaker values for voice cloning. Defaults to `get_user_data_dir()`.

 v: dev ▼

# BarkArgs

[↑ Back to top](#)

## Bark Model

```
class TTS.tts.models.bark.Bark(config, tokenizer=BertTokenizer(name_or_path='bert-base-multilingual-cased', vocab_size=119547, model_max_length=512, is_fast=False, padding_side='right', truncation_side='right', special_tokens={'unk_token': '[UNK]'),
```

 v: dev ▼

```
'sep_token': '[SEP]', 'pad_token': '[PAD]', 'cls_token': '[CLS]', 'mask_token': '[MASK]'}, clean_up_tokenization_spaces=True))
```

[\[source\]](#)

```
generate_audio(text, history_prompt, temp=0.7, waveform_temp=0.7, base=None, allow_early_stop=True, **kwargs,
```

↑ Back to top

[\[source\]](#)

Generate audio array from input text.

#### PARAMETERS:

- **text** – text to be turned into audio
- **history\_prompt** – history choice for audio cloning
- **text\_temp** – generation temperature (1.0 more diverse, 0.0 more conservative)
- **waveform\_temp** – generation temperature (1.0 more diverse, 0.0 more conservative)

#### RETURNS:

numpy audio array at sample frequency 24khz

```
generate_voice(audio, speaker_id, voice_dir)
```

[\[source\]](#)

Generate a voice from the given audio and text.

#### PARAMETERS:

- **audio** (*str*) – Path to the audio file.
- **speaker\_id** (*str*) – Speaker name.
- **voice\_dir** (*str*) – Path to the directory to save the generate voice.

```
load_checkpoint(config, checkpoint_dir, text_model_path=None, coarse_model_path=None, fine_model_path=None, eval=False, strict=True, **kwargs)
```

[\[source\]](#)

Load a model checkpoints from a directory. This model is with multiple checkpoint files and it expects to have all the files to be under the given *checkpoint\_dir* with the right names. If

eval is True, set the model to eval mode.

#### PARAMETERS:

- **config** ([TortoiseConfig](#)) – The ↑ Back to top
- **checkpoint\_dir** (*str*) – The directory where the checkpoints are stored.
- **ar\_checkpoint\_path** (*str, optional*) – The path to the autoregressive checkpoint. Defaults to None.
- **diff\_checkpoint\_path** (*str, optional*) – The path to the diffusion checkpoint. Defaults to None.
- **clvp\_checkpoint\_path** (*str, optional*) – The path to the CLVP checkpoint. Defaults to None.
- **vocoder\_checkpoint\_path** (*str, optional*) – The path to the vocoder checkpoint. Defaults to None.
- **eval** (*bool, optional*) – Whether to set the model to eval mode. Defaults to False.
- **strict** (*bool, optional*) – Whether to load the model strictly. Defaults to True.

**semantic\_to\_waveform**(semantic\_tokens, history\_prompt=None, temp=0.7, base=None)[\[source\]](#)

Generate audio array from semantic input.

#### PARAMETERS:

- **semantic\_tokens** – semantic token output from *text\_to\_semantic*
- **history\_prompt** – history choice for audio cloning
- **temp** – generation temperature (1.0 more diverse, 0.0 more conservative)

#### RETURNS:

numpy audio array at sample frequency 24khz

**synthesize**(text, config, speaker\_id='random', voice\_dirs=None, \*\*kwargs) [\[source\]](#)



Synthesize speech with the given input text.

#### PARAMETERS:

- **text** (*str*) – Input text.
- **config** (*BarkConfig*) – Config with inference parameters.
- **speaker\_id** (*str*) – One of the available speaker names. If *random*, it generates a random speaker.
- **speaker\_wav** (*str*) – Path to the speaker audio file for cloning a new voice. It is cloned and saved in *voice\_dirs* with the name *speaker\_id*. Defaults to None.
- **voice\_dirs** (*List[str]*) – List of paths that host reference audio files for speakers. Defaults to None.
- **\*\*kwargs** – Model specific inference settings used by *generate\_audio()* and `TTS.tts.layers.bark.inference_funcs.generate_text_semantic()`.

↑ Back to top

#### RETURNS:

A dictionary of the output values with *wav* as output waveform, *deterministic\_seed* as seed used at inference, *text\_input* as text token IDs after tokenizer, *voice\_samples* as samples used for cloning, *conditioning\_latents* as latents used at inference.

```
text_to_semantic(text, history_prompt=None, temp=0.7, base=None,
                  allow_early_stop=True, **kwargs)
```

[\[source\]](#)

Generate semantic array from text.

#### PARAMETERS:

- **text** – text to be turned into audio
- **history\_prompt** – history choice for audio cloning
- **temp** – generation temperature (1.0 more diverse, 0.0 more conservative)

#### RETURNS:

numpy semantic array to be fed into *semantic\_to\_waveform*



Copyright © 2021 Coqui GmbH, 2020 TTS authors  
Made with [Sphinx](#) and @pradyunsg's [Furo](#)

v: dev ▼